

Bharat Evidence Agent

From authorized evidence to reviewable action.

- Local, offline triage for exported messages and payments.
- Understand > Reason > Plan > Use Tools > Act.
- Synthetic demo. Human-reviewed findings. No device bypass.

Make review leads traceable

A practical workflow for mixed-language evidence exports.

- Messages and payment records need a shared timeline.
- English / Hindi patterns support a Bharat-focused UPI demo.
- Each lead must cite its source and explain its limits.
- Target users: authorized analysts and citizen-support teams.
- Impact is a hypothesis; no field adoption or time savings claimed.

Plan. Inspect. Replan. Act.

A deterministic tool agent—not a generic chatbot.

- Intake: validate, normalize, compute import hashes, save case.
- Plan: select integrity, timeline, entity and scan tools by objective.
- Reason: if flagged messages + payments exist, add correlation.
- Use tools: trace evidence IDs and explain the 30-minute window.
- Act: persist review packet; human approval unlocks final report.

Inspectable from input to report

Python + SQLite + browser UI. No runtime API key or dependencies.

- Synthetic demo: 7 records > 3 review leads + 6 timed events.
- 20 automated tests cover tools, validation, audit and API controls.
- Goal change to timeline skips fraud scan and payment correlation.
- Hashes protect import snapshots—not acquisition authenticity.
- Correlation means proximity, not causation or proof of fraud.

Built locally. Reviewed by humans.

A focused prototype with a clear production boundary.

- Current scope: local single-user, 5 MB / 5,000-record imports.
- Roadmap: reviewer authentication, encryption, signed custody.
- Next: validate detectors and UX with authorized practitioners.
- Team, GitHub and demo URL: owner must add real details.
- Official PDF missing: full rules and deadline compliance unverified.